

Security properties

- Confidentiality (C)
 - Data secrecy: Keep data secret from unauthorized subjects
 - Privacy: Individuals control who can access their information
- Integrity (I)
 - Aka “authenticity”
 - Data integrity: Keep data from being modified
 - System integrity: Keep systems functioning as intended
- Availability (A)
 - Keep the system running and responsive to legitimate clients

cryptography

- Authentication
 - verify that sender sent the message + not changed (digital signature)
- Integrity
 - message has not been modified
- Non-repudiation
 - sender cannot deny that she indeed sent the message

functions

- One-way hash functions $h = H(M)$: no key
 - for integrity
 - Plaintext \rightarrow fingerprint
 - Brute force attack \rightarrow choose a large output size 128-160
 - Take a variable-length input M and produce fixed-length output
 - Public
 - weak collision resistance: Given M it is very hard to find M' such that $H(M) = H(M')$
 - cheating after send: once the message is sent, sender cannot change it and claim they sent a different message
 - strong collision resistance: hard to find two random messages M_1 and M_2 such that $H(M_1) = H(M_2)$
 - cheating before send: sender cannot prepare two messages that have the same hash, send one and claim they have sent the other
- Symmetric crypto: one key
 - for confidentiality
 - only Alice and Bob know the same *shared key*, sender will encrypt with that key and receiver will decrypt
 - Plaintext \rightarrow ciphertext
 - Substitution
 - Monoalphabetic – each character is replaced with another character
 - Homophonic – each character is replaced with a character chosen randomly from a subset

- Polygram – each sequence of characters of length n is replaced with another sequence of characters of length n
- Polyalphabetic – many monoalphabetic ciphers are used sequentially, stream ciphers: a bit or a byte at a time
 - Assume XOR (weak but fast) with the key (long and random)
 - $A \text{ xor } K = B$
 - $A \text{ xor } B = K$
- One-time pad - Polyalphabetic cipher with **infinite** key, **randomly** generated, perfectly synchronized
- Block ciphers: polygram, block size = key size
 - Electronic Code Book (ECB) No dependency, do in parallel
 - Mallory can detect mapping, replay, fabricate, add/drop/replace and not be detected
 - Cipher Block Chaining (CBC) Dependency on earlier blocks
 - same plaintext blocks will encrypt to different ciphertext blocks, encryption/decryption cannot be parallelized
 - Replay is still possible if we don't use timestamps
 - Initialization vector (IV)
- Transposition(shuffling)
- Asymmetric crypto: two keys

$$D_{K2}(E_{K1}(M)) = M$$
 - for confidentiality and authenticity
 - Alice has *public key* and *private key*
 - Everyone knows Alice's public key but only Alice knows her private key
 - Plaintext -> ciphertext
 - Confidentiality, A use B's public key to encrypt, B use its private key to decrypt
 - Authenticity, A use A's private key to sign, anyone use A's public key to verify
 - Functionality is greater but much slower
 - RSA
 - modular exponentiation in Galois Field GF(n) is efficient
 - factor large number is hard for attackers
 - digital signature $E_{\text{privA}}(H(M))$ or $E_{\text{privA}}(M)$, provide non-repudiation

Attacks for symmetric/asymmetric crypto

- Ciphertext-only attack: gather and analyze enough **ciphertext** to learn decryption key, recognize the plaintext
- known-plaintext attack: observe many **ciphertexts** for **known plaintexts** to learn decryption key
- chosen-plaintext attack: feed chosen messages M into encryption algorithm and look at resulting ciphertexts C . Learn decryption key/messages M that produce C

- Man-in-the-middle attack: substitute, modify, drop, replay messages
- Brute-force attack: caught a ciphertext and try every possible key to decrypt -> choosing a large keyspace

Shared Key Exchange

- Diffie-Hellman
 - A sends $g^a \bmod n$
 - B sends $g^b \bmod n$
 - Shared key is **$g^{ab} \bmod n$**
 - Hard to guess a and b with a large n
 - Man in the middle, no authentication
- KDC (Using a trusted third party)
 - Secrets should never be sent in clear, preconfigured keys $K_{KDC,C}$ and $K_{KDC,S}$
 - Use nonces to prevent replay attacks
 - Challenge response 3 flavors: check if S knows the same shared key $K_{C,S}$
 - We should prevent reuse of old keys -> tickets expire after some time, **validity period** in ticket
- Kerberos
 - Authentication server (AS) authenticates users, issues a ticket for clients to talk to TGS
 - Ticket Granting Server (TGS) Issue tickets for clients to talk to servers
 - **TGS+AS = KDC**
 - Each ticket has a validity period: timestamp and lifetime
 - Each service request (client to server) has an authenticator (nonce): timestamp + client identity, encrypted with a session key
- public key cryptography
 - A sends $E_{Pub_B}(K_{AB})$ to B, much slower

Public Key Exchange

- Need a trusted third party – Trent, aka *Certificate Authority (CA)*
- Everyone knows Trent's public key
- Trent signs Public-key Certificate: $E_{privateKeyTrent}(\text{public key, Id(dns name)})$
- Certificate-Based Key Exchange
- Recovery from Stolen Private Keys

Password Authentication

- Dictionary attack – guessing passwords
 - Offline: **comparing** guesses to a list of precomputed hashes of **popular** pass. -> use a random number salt with a password in hashing
 - Online: trying popular passwords **manually** on server UI -> hard because password table is stored in the file `/etc/shadow`, only accessible by superuser
- Personal information attack

- Online: Try to guess a specific user's password by using **personal information** about the user manually on server UI
- Reuse attack
 - Steal a password from one server, try it at other servers
- Lamport hash or S-KEY – time-varying password
 - Someone **sniffing** on the network can learn the password if it is transmitted in clear
 - Encrypt communication
 - If **replay** attacks are possible, someone can still steal passwords and reuse them
 - Host sets password $x_0=h(R)$, $x_1=h(h(R))$, $x_2=h(h(h(R)))$, ..., x_{100}
 - User log on with x_{100} , x_{99} , ...
- Shared-Key Authentication challenge response
- Public-Key Authentication
- Single Sign-On SSO
- One Identity Provider
 - Various systems share trust with **one** identity provider
 - Identity provider returns an encrypted token to the system, using a shared key
- Multiple Identity Providers
 - Various systems share trust with **all** identity providers
 - System decides which identity provider (IdP) to use, redirects user to identity provider
- Cookie-based Authentication
 - Placed into browser cache by servers
 - persistent sign in, shopping cart, user preferences
- Token-based Authentication
- Biometric Authentication

Access Control

- Discretionary (owner grant access)
 - Access Control Matrix
 - Access Control Lists (ACLs) For each object
 - Capabilities For each principal
- Mandatory (by system, government/army)
 - Each object has a **classification** and each subject has a **clearance**
 - **top secret > secret > confidential > restricted > unclassified**
 - **no read-up**
 - **no append-down**
 - Append is allowed on objects of the **same or higher** classification
 - Write is allowed **only** on objects of the **same** classification
 - Trusted subjects – the “no write-down” rule does not apply to them
- Bell-LaPadula Policy Model

- Combine mandatory and discretionary
- Intersection of Access Control Matrix and level check
- Role-based
 - depends on one's role in the organization
 - Maps to organization structure(companies)
- Attribute-based
 - depends on *attributes* assigned to user and object, and on *environment attributes*